

Loan_Eligibility_Predictions

November 26, 2023

1 Extract data from the source

```
[39]: import pandas as pd
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import ( confusion_matrix, accuracy_score,
    precision_score, recall_score, f1_score)
```

```
[2]: loan_predictions = pd.read_csv('Loan_Eligibility_Predictions.csv')
```

```
[3]: loan_predictions.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 614 entries, 0 to 613
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Loan_ID               614 non-null   object
1   Gender                601 non-null   object
2   Married               611 non-null   object
3   Dependents            599 non-null   object
4   Education              614 non-null   object
5   Self_Employed         582 non-null   object
6   ApplicantIncome       614 non-null   int64
7   CoapplicantIncome     614 non-null   float64
8   LoanAmount            592 non-null   float64
9   Loan_Amount_Term      600 non-null   float64
10  Credit_History        564 non-null   float64
11  Property_Area         614 non-null   object
12  Loan_Status           614 non-null   object
dtypes: float64(4), int64(1), object(8)
memory usage: 62.5+ KB
```

```
[4]: loan_predictions.head()
```

```
[4]:      Loan_ID Gender Married Dependents      Education Self_Employed \
0  LP001002   Male      No           0      Graduate         No
1  LP001003   Male     Yes           1      Graduate         No
2  LP001005   Male     Yes           0      Graduate         Yes
3  LP001006   Male     Yes           0  Not Graduate         No
4  LP001008   Male      No           0      Graduate         No

      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
0              5849              0.0          NaN          360.0
1              4583             1508.0          128.0          360.0
2              3000              0.0           66.0          360.0
3              2583             2358.0          120.0          360.0
4              6000              0.0          141.0          360.0

      Credit_History Property_Area Loan_Status
0              1.0         Urban           Y
1              1.0         Rural           N
2              1.0         Urban           Y
3              1.0         Urban           Y
4              1.0         Urban           Y
```

2 Exploratory Data Analysis (EDA)

2.0.1 Identify descriptive statistics on variables/features

```
[5]: # Descriptive statistics for numerical columns
loan_predictions.describe()
```

```
[5]:      ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term \
count      614.000000      614.000000  592.000000      600.000000
mean      5403.459283      1621.245798  146.412162      342.000000
std       6109.041673      2926.248369   85.587325       65.12041
min       150.000000        0.000000    9.000000       12.000000
25%       2877.500000        0.000000  100.000000      360.000000
50%       3812.500000      1188.500000  128.000000      360.000000
75%       5795.000000      2297.250000  168.000000      360.000000
max      81000.000000     41667.000000  700.000000      480.000000

      Credit_History
count      564.000000
mean        0.842199
std         0.364878
min         0.000000
25%         1.000000
50%         1.000000
75%         1.000000
max         1.000000
```

```
[6]: # Count for categorical columns
categorical_columns = ['Gender', 'Married', 'Dependents', 'Education', '
↳ 'Self_Employed', 'Credit_History', 'Property_Area', 'Loan_Status']

for column in categorical_columns:
    print(loan_predictions[column].value_counts())
    print("\n")
```

Gender

Male 489

Female 112

Name: count, dtype: int64

Married

Yes 398

No 213

Name: count, dtype: int64

Dependents

0 345

1 102

2 101

3+ 51

Name: count, dtype: int64

Education

Graduate 480

Not Graduate 134

Name: count, dtype: int64

Self_Employed

No 500

Yes 82

Name: count, dtype: int64

Credit_History

1.0 475

0.0 89

Name: count, dtype: int64

Property_Area

Semiurban 233

```
Urban      202
Rural      179
Name: count, dtype: int64
```

```
Loan_Status
Y      422
N      192
Name: count, dtype: int64
```

```
[7]: loan_predictions.select_dtypes(include=['object']).describe()
```

```
[7]:      Loan_ID  Gender  Married  Dependents  Education  Self_Employed  \
count      614      601       611         599         614         582
unique      614         2         2           4           2           2
top    LP001002    Male     Yes           0  Graduate         No
freq           1      489      398        345         480        500
```

```
      Property_Area  Loan_Status
count           614           614
unique            3             2
top      Semiurban           Y
freq           233           422
```

2.0.2 Determine correlation between variables/features

```
[8]: # Determine correlation between variables/features (R2)
numeric_columns = loan_predictions.select_dtypes(include=[np.number])
comatrix = numeric_columns.corr()
comatrix
```

```
[8]:      ApplicantIncome  CoapplicantIncome  LoanAmount  \
ApplicantIncome      1.000000      -0.116605      0.570909
CoapplicantIncome    -0.116605      1.000000      0.188619
LoanAmount           0.570909      0.188619      1.000000
Loan_Amount_Term     -0.045306     -0.059878      0.039447
Credit_History       -0.014715     -0.002056     -0.008433
```

```
      Loan_Amount_Term  Credit_History
ApplicantIncome      -0.045306     -0.014715
CoapplicantIncome    -0.059878     -0.002056
LoanAmount           0.039447     -0.008433
Loan_Amount_Term      1.000000      0.001470
Credit_History       0.001470      1.000000
```

2.0.3 Identify and handle NULL values

```
[9]: # Check for missing values in the entire DataFrame
loan_predictions.isnull().sum()
```

```
[9]: Loan_ID          0
      Gender          13
      Married         3
      Dependents     15
      Education       0
      Self_Employed  32
      ApplicantIncome 0
      CoapplicantIncome 0
      LoanAmount     22
      Loan_Amount_Term 14
      Credit_History  50
      Property_Area   0
      Loan_Status     0
      dtype: int64
```

```
[10]: loan_predictions.fillna({'Gender': 'Male', 'Married': 'Yes', 'Dependents': '0',
    ↪ 'Self_Employed': 'No', 'LoanAmount': loan_predictions['LoanAmount'].
    ↪ median(), 'Loan_Amount_Term': loan_predictions['Loan_Amount_Term'].
    ↪ mode()[0], 'Credit_History': 1}, inplace=True)
loan_predictions.isnull().sum()
```

```
[10]: Loan_ID          0
      Gender          0
      Married         0
      Dependents     0
      Education       0
      Self_Employed  0
      ApplicantIncome 0
      CoapplicantIncome 0
      LoanAmount     0
      Loan_Amount_Term 0
      Credit_History  0
      Property_Area   0
      Loan_Status     0
      dtype: int64
```

2.0.4 Identify any outliers

```
[11]: # Filter to select only numeric columns
numeric_columns = loan_predictions.select_dtypes(include=np.number)

# Calculate IQR for each numeric column
Q1 = numeric_columns.quantile(0.25)
```

```

Q3 = numeric_columns.quantile(0.75)
IQR = Q3 - Q1

# Identify potential outliers using the IQR
outliers = (numeric_columns < (Q1 - 1.5 * IQR)) | (numeric_columns > (Q3 + 1.5 * IQR))
outliers.sum()

```

```

[11]: ApplicantIncome      50
      CoapplicantIncome    18
      LoanAmount           41
      Loan_Amount_Term     88
      Credit_History       89
      dtype: int64

```

2.0.5 Handling outliers

```

[12]: # Define the lower and upper bounds for outliers
      lower_bound = Q1 - 1.5 * IQR
      upper_bound = Q3 + 1.5 * IQR

      # Replace outliers with the nearest non-outlier value
      for column in numeric_columns.columns:
          numeric_columns.loc[numeric_columns[column] < lower_bound[column], column] = lower_bound[column]
          numeric_columns.loc[numeric_columns[column] > upper_bound[column], column] = upper_bound[column]

```

```

[13]: outliers_after_handling = (numeric_columns < (Q1 - 1.5 * IQR)) | (numeric_columns > (Q3 + 1.5 * IQR))
      outliers_count = outliers_after_handling.sum()
      outliers_count

```

```

[13]: ApplicantIncome      0
      CoapplicantIncome    0
      LoanAmount           0
      Loan_Amount_Term     0
      Credit_History       0
      dtype: int64

```

3 Visualizations


3.0.1 Histograms of relevant variables

```
[14]: # Histograms of applicants income and loan amount
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

sns.histplot(loan_predictions['ApplicantIncome'], bins=20, kde=True, ax=axes[0])
axes[0].set_title('Applicant\'s Income Distribution')

sns.histplot(loan_predictions['LoanAmount'], bins=20, kde=True, ax=axes[1])
axes[1].set_title('Loan Amount Distribution')

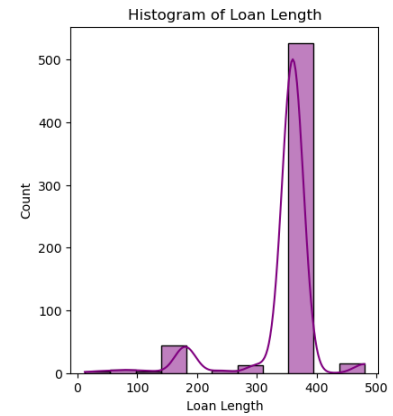

plt.tight_layout()
plt.show()
```



```
[15]: # Histogram of credit history and loan length
plt.figure(figsize=(15, 5))
plt.subplot(1, 3, 1)
sns.histplot(data=loan_predictions, x='Credit_History', kde=True, color='red')
plt.title("Histogram of Credit History")
plt.xlabel("Credit History")

plt.subplot(1, 3, 3)
sns.histplot(data=loan_predictions, x='Loan_Amount_Term', kde=True,
             color='purple')
plt.title("Histogram of Loan Length")
plt.xlabel("Loan Length")
```


```
[15]: Text(0.5, 0, 'Loan Length')
```



3.0.2 Different types of plots

Loan Status with respect to Gender

```
[16]: # Count plot for 'Loan_Status' with respect to 'Gender'
sns.countplot(data=loan_predictions, x='Loan_Status', hue='Gender')
plt.title('Count of Loan Status based on Gender')
plt.xlabel('Loan Status')
plt.ylabel('Count')
plt.show()
```


Approval and Denial rates based on credit history

```
[17]: # Calculate approval and denial counts based on credit history
credit_history_counts = loan_predictions.groupby(['Credit_History',
↳ 'Loan_Status']).size().unstack()

# Calculate the total counts of approvals and denials
total_counts = credit_history_counts.sum(axis=1)

# Plotting the donut chart
fig, ax = plt.subplots()
ax.pie(total_counts, labels=total_counts.index, autopct='%1.1f%%',
↳ startangle=90, colors=['lightblue', 'lightgreen'])
centre_circle = plt.Circle((0,0),0.70,fc='white')
fig = plt.gcf()
fig.gca().add_artist(centre_circle)
ax.axis('equal')
plt.title('Approval and Denial Rates by Credit History')
plt.show()
```

Approval and Denial Rates by Credit History



What percentage of approved applications have a good credit history?


```
[18]: # Calculate the number of approved applications
approved_applications = loan_predictions[loan_predictions['Loan_Status'] == 'Y']

# Calculate the percentage of approved applications with a good credit history
good_credit_approved = □
    ↳ approved_applications[approved_applications['Credit_History'] == 1]
percentage_good_credit_approved = (len(good_credit_approved) / □
    ↳ len(approved_applications)) * 100

# Calculate the percentage of approved applications with a bad credit history
percentage_bad_credit_approved = 100 - percentage_good_credit_approved

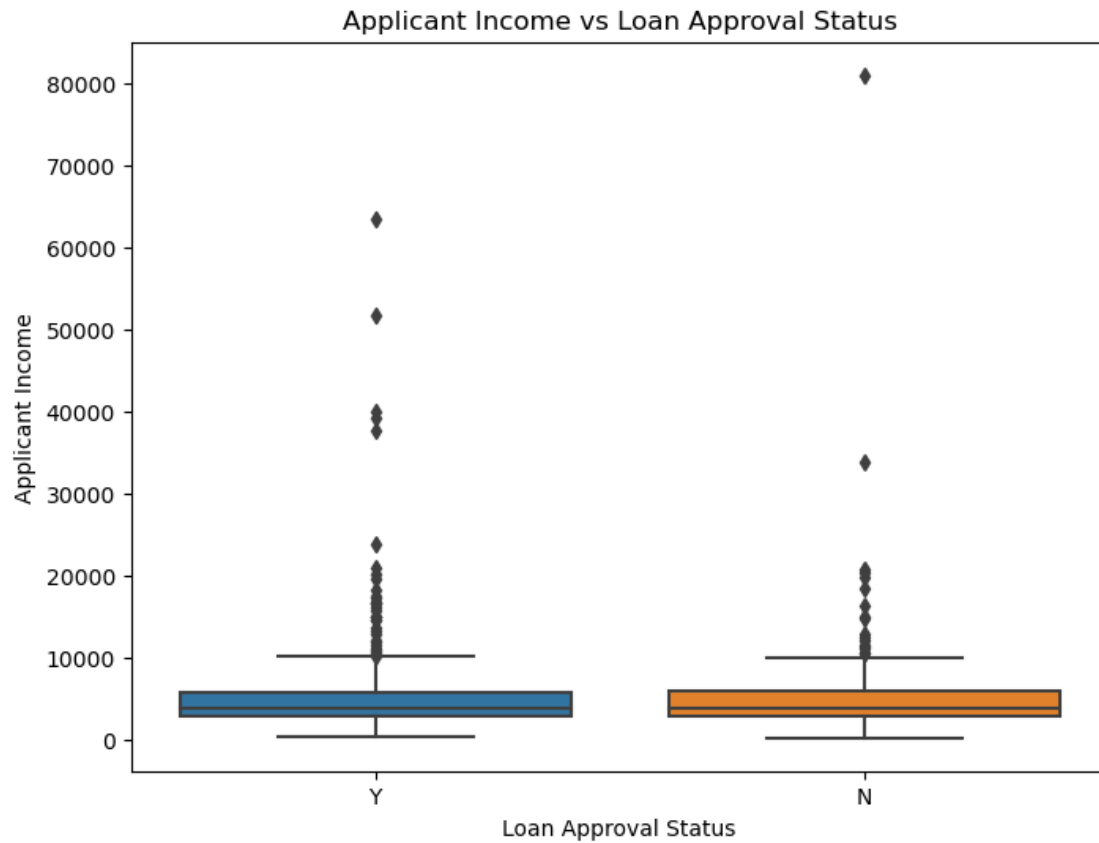
# Plotting
plt.pie([percentage_good_credit_approved, percentage_bad_credit_approved], □
    ↳ labels=['Good Credit History', 'Bad Credit History'], colors=['lightblue', □
    ↳ 'lightcoral'], autopct='%1.1f%%')
plt.axis('equal')
plt.title('Percentage of Approved Applications with Good Credit History')
plt.show()
```

Percentage of Approved Applications with Good Credit History



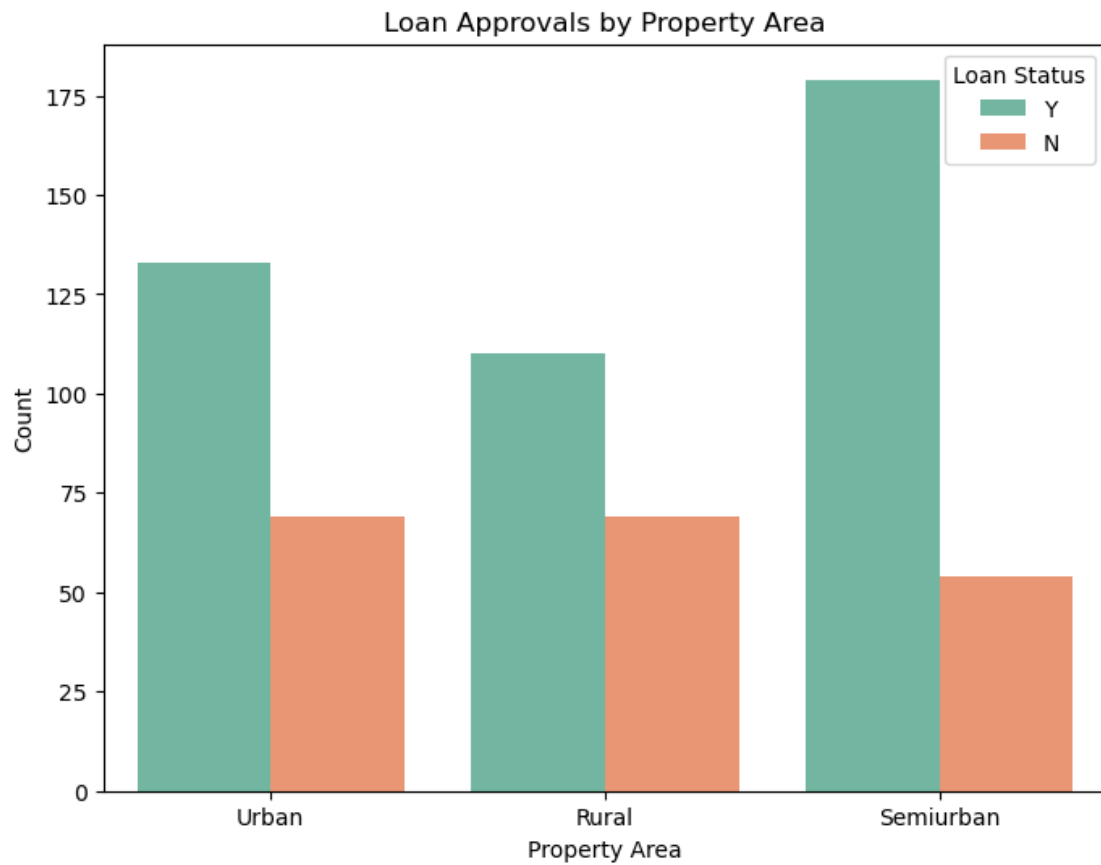
Are there income thresholds that significantly impact loan approval?

```
[19]: # Create a boxplot of ApplicantIncome for approved and denied loans
plt.figure(figsize=(8, 6))
sns.boxplot(x='Loan_Status', y='ApplicantIncome', data=loan_predictions)
plt.title('Applicant Income vs Loan Approval Status')
plt.xlabel('Loan Approval Status')
plt.ylabel('Applicant Income')
plt.show()
```



Loan approval by property area


```
[20]: # Plotting loan approvals by Property_Area
plt.figure(figsize=(8, 6))
sns.countplot(x='Property_Area', hue='Loan_Status', data=loan_predictions,
              palette='Set2')
plt.title('Loan Approvals by Property Area')
plt.xlabel('Property Area')
plt.ylabel('Count')
plt.legend(title='Loan Status', loc='upper right')
plt.show()
```



Correlations between variables

```
[33]: # Plot the heatmap
plt.figure(figsize=(12, 8))
sns.heatmap(comatrix, annot=True)
```

[33]: <Axes: >



4 Model Development using KNN

```
[34]: X = loan_predictions[['Credit_History', 'ApplicantIncome', 'CoapplicantIncome',
↪ 'LoanAmount', 'Loan_Amount_Term']]
y = loan_predictions['Loan_Status']

[35]: y = y.map({'N': 0, 'Y': 1})

[36]: # Step 1: Split the data into training and test datasets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↪ random_state=42)

[37]: # Step 2: Standardize data to the same scale
standardizer = StandardScaler()
x_standardized = standardizer.fit_transform(X_train)
x_test_std = standardizer.fit_transform(X_test)

[40]: # Step 3: Create model (fit the training data)
knn = KNeighborsClassifier(n_neighbors=17).fit(x_standardized, y_train)

[41]: # Step 4: Use model to predict the class of test data
y_predicted= knn.predict(x_test_std)
```


```
[42]: # Step 5: Model Evaluation
print(accuracy_score(y_test, y_predicted))
print(precision_score(y_test, y_predicted, average= 'macro'))
print(recall_score(y_test, y_predicted, average= 'macro'))
print(f1_score(y_test, y_predicted, average= 'macro'))
```

```
0.7967479674796748
0.8809523809523809
0.7093023255813954
0.7275143996455471
```

```
[43]: # Step 6: Confusion Matrix
conf_matrix = confusion_matrix(y_test, y_predicted)
conf_matrix
```

```
[43]: array([[18, 25],
          [ 0, 80]], dtype=int64)
```


```
[44]: # Plot confusion matrix
plt.figure(figsize=(8,8))
sns.set(font_scale = 1.5)
ax = sns.heatmap(
    conf_matrix,
    annot=True,
    fmt='d',
    cbar=False,
    cmap='flag',
    vmax=175
)
ax.set_xlabel("Predicted", labelpad=20)
ax.set_ylabel("Actual", labelpad=20)
plt.show()
```



```
[45]: plot_data= {'k' : [9,25,3,5,7], 'f1_score': [.71, .75, .68, .70, .72]}
      sns.relplot(data=plot_data, kind='line', x='k', y='f1_score')
```

C:\ProgramData\anaconda3\Lib\site-packages\seaborn\axisgrid.py:118: UserWarning:
The figure layout has changed to tight
self._figure.tight_layout(*args, **kwargs)

```
[45]: <seaborn.axisgrid.FacetGrid at 0x1c6c036a590>
```

```
[ ]: # Bend is observed at 9, Optimal K is 9
```